

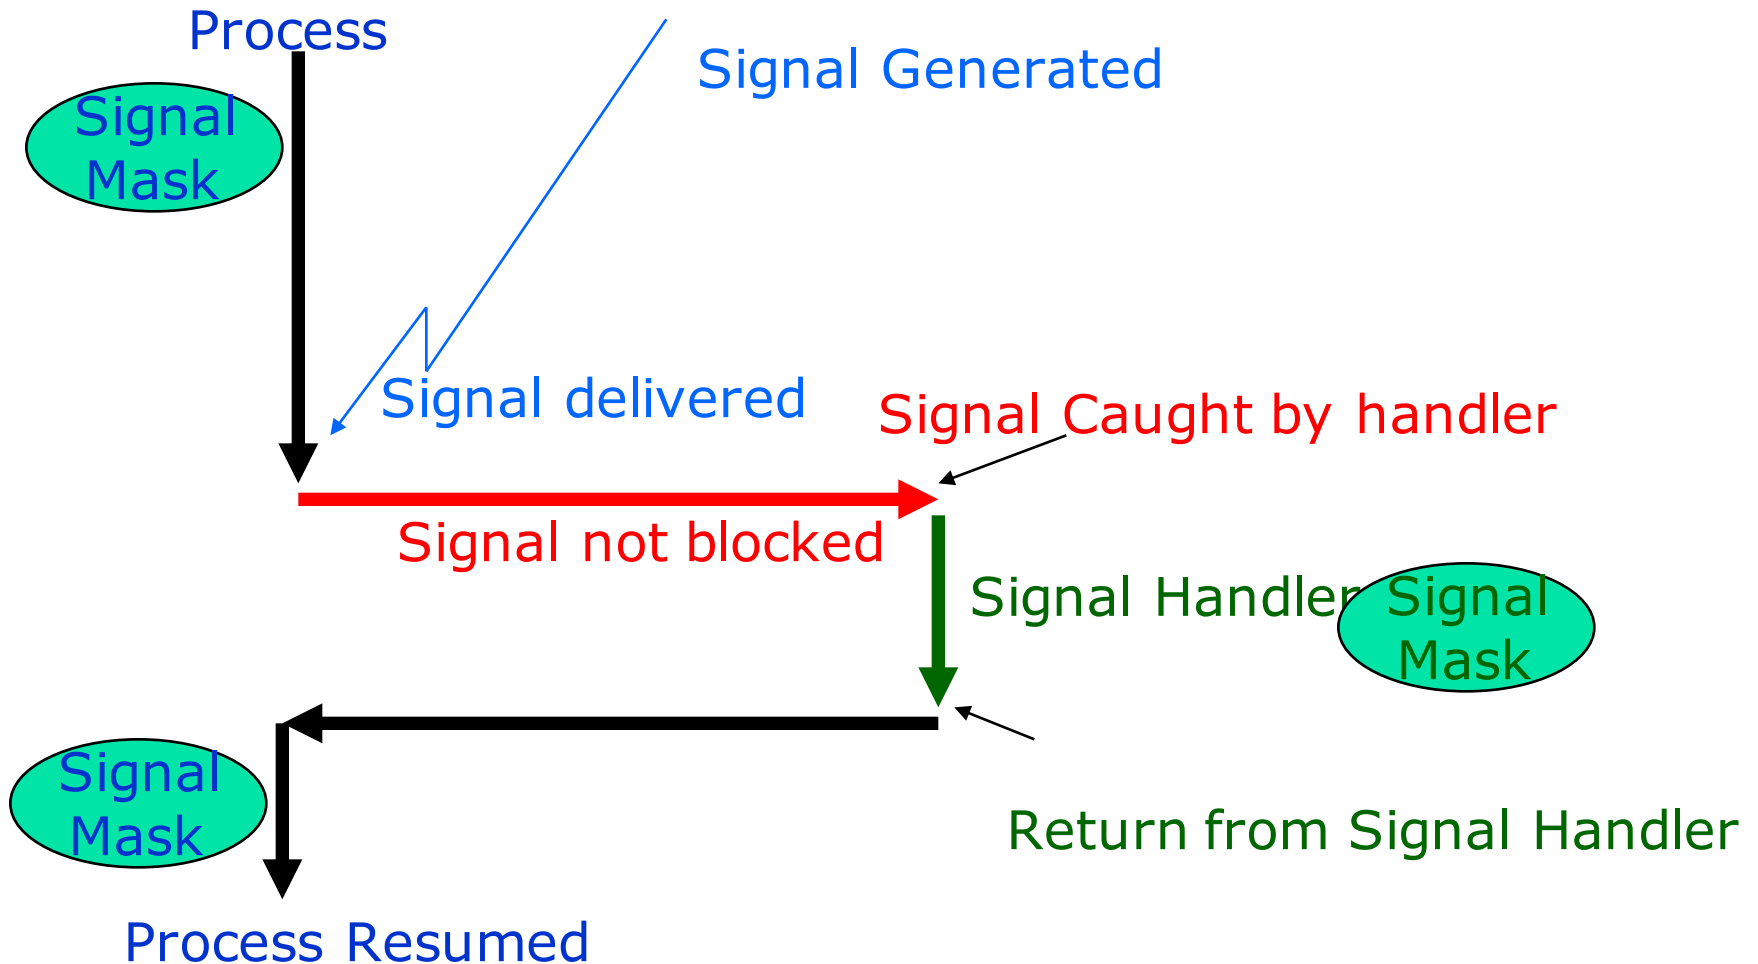
Introduction to signals

- Concept
- POSIX signals
- Generating signals
- Catching signals
- Signal sets
- Signal masks
- Waiting for signals

Signals

- A signal is a software notification to a process of an event.
- Needed to enable asynchronous events
- Examples of asynchronous events:
 - **Email message arrives** on my machine – mailing agent (user) process should retrieve
 - **Invalid memory access happens** – OS should inform scheduler to remove process from the processor
 - **Alarm clock goes on** – process which sets the alarm should catch it

How Signals Work



Examples of POSIX Signals

Signal	Description	default action
SIGABRT	process abort	implementation dependent
SIGALRM	alarm clock	abnormal termination
SIGBUS	access undefined part of memory object	implementation dependent
SIGCHLD	child terminated, stopped or continued	ignore
SIGILL	invalid hardware instruction	implementation dependent
SIGINT	interactive attention signal (usually ctrl-C)	abnormal termination
SIGKILL	terminated (cannot be caught or ignored)	abnormal termination

Examples of POSIX Signals

Signal	Description	default action
SIGSEGV	Invalid memory reference	implementation dependent
SIGSTOP	Execution stopped	stop
SIGTERM	termination	Abnormal termination
SIGTSTP	Terminal stop	stop
SIGTTIN	Background process attempting read	stop
SIGTTOU	Background process attempting write	stop
SIGURG	High bandwidth data available on socket	ignore
SIGUSR1	User-defined signal 1	abnormal termination

Generating Signals

- Signal has a symbolic name starting with SIG
- Signal names are defined in signal.h
- Users can generate signals (e.g., SIGUSR1)
- OS generates signals when certain errors occur (e.g., SIGSEGV – invalid memory reference)
- Specific calls generate signals such as alarm (e.g., SIGALRM)

Command Line Generates Signals

- You can send a signal to a process from the command line using **kill**
- `kill -l` will list the signals the system understands
- `kill [-signal] pid` will send a signal to a process.
 - The optional argument may be a name or a number (default is SIGTERM).
- To unconditionally kill a process, use:
 - `kill -9 pid` which is
`kill -SIGKILL pid`.

Command Line Generates Signals

- CTRL-C is SIGINT (interactive attention signal)
- CTRL-Z is SIGSTOP (execution stopped – cannot be ignored)
- CTRL-Y is SIGCONT (execution continued if stopped)
- CTRL-D is SIGQUIT (interactive termination: core dump)

Signal Masks

- Process can temporarily prevent signal from being delivered by *blocking* it.
- *Signal Mask* contains a set of signals currently blocked.
- **Important!** Blocking a signal is different from ignoring signal. Why?
- When a process blocks a signal, the OS does not deliver signal until the process unblocks the signal
 - A *blocked* signal is not delivered to a process until it is unblocked.
- When a process ignores signal, signal is delivered and the process handles it by throwing it away.

Signal Masks

SigMask

SigInt	SigQuit	SigKill	...	SigCont	SigAbrt
--------	---------	---------	-----	---------	---------

0	0	1	...	1	0
---	---	---	-----	---	---

Signal SigInt Bit 2, Signal Sigkill Bit 9,
Signal SigChld Bit 20

A SIGSET is a collection of signals:
#000003 is SIGHUP + SIGINT

Manipulate signal mask

- Manipulate signal mask of a process
 - `int sigprocmask(int how, const sigset_t *set, sigset_t *oset);`
 - *how*: SIG_BLOCK, SIG_UNBLOCK, SIG_SETMASK
 - *set*: New set
 - *oset*: Old set
 - See `example2.c` and `example3.c`

Manipulate signal sets

- `#include <signal.h>`

`int sigemptyset(sigset_t *set);`

`int sigfillset(sigset_t *set);`

`int sigaddset(sigset_t *set, int signo);`

`int sigdelset(sigset_t *set, int signo);`

`int sigismember(const sigset_t *set, int signo);`

Example: Blocking signals in a critical region

- For a critical region where you don't want certain signal to come, the program will look like:

```
sigprocmask(SIG_BLOCK, &newmask, &oldmask);  
..... /* critical region */  
sigprocmask(SIG_SETMASK, &oldmask, NULL);
```

Catching and Ignoring Signals - SIGACTION

- sigaction function allows the caller to examine or specify action associated with a specific signal
- Program installs *signal handler* by calling sigaction with the name of a user-written function.
- The function sigaction is used to specify what is to happen to a signal when it is delivered.

sigaction

– Supersedes the signal function

- `#include <signal.h>`
- `int sigaction(int signo, const struct sigaction * act, struct sigaction *oact)`
- `struct sigaction`
 - `{`
 - `void (*sa_handler)(); /* signal handler */`
 - `sigset_t sa_mask; /*additional signal to block */`
 - `int sa_flags; /* options: we will set it to 0 */`
 - `sa_sigaction; /* we will not use it */`
 - `};`

– See example4.c

kill

- Send a signal to a process
 - `#include <signal.h>`
 - `#include <sys/types.h>`
 - `int kill(pid_t pid, int signo);`
 - See `example5.c`

Timers Generate SIGALRM Signals

- `#include <unistd.h>`
- `unsigned alarm (unsigned seconds);`
- `alarm(20)` creates SIGALRM to calling process after 20 real time seconds.
- Calls are not stacked
- `alarm(0)` cancels alarm
- See `example6.c` for the use of *alarm*

Impact of signals on system calls

- A system call may return prematurely
 - See example7.c
- How can we deal with this problem?
 - Check the return value of the system call and act accordingly

Waiting for Signals

- Signals provide method for waiting for event without busy waiting
- Busy waiting
 - Means continually using CPU cycles to test for occurrence of event
 - Means a program does this testing by checking the value of variable in loop
- More Efficient Approach
 - Suspend process until waited-for event occurs
- POSIX pause, sigsuspend, sigwait provide three mechanisms to suspend process until signal occurs
- pause() – too simple, and problematic if signals are delivered before pause (cannot unblock the signal and start pause() in atomic way)

sigsuspend

- sigsuspend function sets signal mask and **suspends process until signal is caught** by the process
- sigsuspend returns when signal handler of the caught signal returns
- `#include <signal.h>`
- `int sigsuspend(const sigset_t *sigmask);`

Example: sigsuspend

- What's wrong?

```
sigfillset(&sigmask);  
sigdelset(&sigmask, signal);  
sigsuspend(&sigmask);
```

- signal is the only **signal unblocked** that can cause sigsuspend to return
- If the signal signal is delivered before the start of the code segment, the process still suspends itself and **deadlocks** if another signal is not generated

Example: sigsuspend (Correct Way to Wait for Single Signal)

```
static volatile sig_atomic_t sigreceived =0;
/*assume signal handler has been setup for signum and it sets sigreceived=1 */

sigset_t maskall, maskmost, maskold;
int signum = SIGUSR1;

sigfillset(&maskall);
sigfillset(&maskmost);
sigdelset(&maskmost, signum);
sigprocmask(SIG_SETMASK, &maskall, &maskold);
if (sigreceived == 0)
    sigsuspend(&maskmost);
sigprocmask(SIG_SETMASK, &maskold, NULL);
```

Use sigwait

- **This function is much cleaner and avoids races and errors !!!!**
- 1. First block all signals
- 2. Put the signals you want to wait for in sigset_t
- 3. Call sigwait
- 4. sigwait blocks the process until at least one of these signals is pending.
- 5. It removes one of the pending signals and gives you the corresponding signal number in the second parameter..
- 6. Do what you want: no signal handler needed.
- 7. It returns 0 on success and -1 on error with errno set.

```
#include <signal.h>
```

```
int sigwait(const sigset_t *restrict sigmask, int  
*restrict signo);
```